


Present

# Julia

## Outline

- Julia - A Fresh Approach to Numerical Computation
- Julia - A *Fast* Approach to Numerical Computation
- Programming Paradigms and Multiple Dispatch
- Hands on Matching Pursuit

## Preface



Everyone is entitled to their personal opinion, workflow and preferences. The focus should be on the scientific progress.

# Julia : Raison d'être and Goals [5]

---

*In short, because we are greedy.*

- The speed of C
- The dynamism of Ruby
- The obvious syntax of Matlab
- The generalizability of Python
- The adhesivity of the Shell
- ...

## Julia - A fresh approach to numerical computation [1]

---

Draft in high level language  $\mapsto$  Reimplement in low level Language

Additional effort:

- Map datastructures correctly
- Ensure composability of datastructures and functions or methods
- Implement interfaces to Open Source Projects
- Add functionality to these
- ...

This is known as the **Two Language Problem** in computer science.

Let's compare Julia with Python ▼

*Python*

- Object Oriented
- Dynamicly Interpreted
- Performant (?)

*Julia*

- Dynamicly Typed
- Just-In-Time Compiled
- Highly performant

# The need for speed

---

f (generic function with 2 methods)

- *# Consider the function*
- `f(x,y) = exp(-(x-y)^2)`

1.1253517471925912e-7

- *# Lets see what happens*
- `f(1, 5)`

```
CodeInfo(
1 - %1 = x - y
   %2 = Core.apply_type(Base.Val, 2)
   %3 = (%2)()
   %4 = Base.literal_pow(Main.workspace285.::^, %1, %3)
   %5 = -%4
   %6 = Main.workspace285.exp(%5)
   return %6
)
```

- *# And under the hood*
- `@code_lowered f(0.2, 3.0)`

[0.379052, 1.40459, 0.589247, -0.778208, -1.24411, 2.10893, -0.696557, -2.43836, 0.260874,

- `begin`
- *# For two arrays*
- `x = fill(1, 100000)`
- `ŷ = randn(100000)`
- `end`

**MethodError: no method matching** `^(::Vector{Float64}, ::Int64)`

Closest candidates are:

```
^(!Matched::Union{AbstractChar, AbstractString}, ::Integer) at strings/basic.jl:718
^(!Matched::Complex{var"#s79"} where var"#s79"<:AbstractFloat, ::Integer) at complex.jl:8
^(!Matched::Complex{var"#s79"} where var"#s79"<:Integer, ::Integer) at complex.jl:820
...
```

```
1. macro expansion @ none:0 [inlined]
2. literal_pow @ none:0 [inlined]
3. f(::Vector{Int64}, ::Vector{Float64}) @ (Other: 2
4. macro expansion @ timing.jl:210 [inlined]
5. (::Main.workspace285.var"#1#2")() @ (Local: 2
6. (::PlutoUI.var"#49#52"{Base.Iterators.Pairs{Union{}, Union{}, Tuple{},
   NamedTuple{(), Tuple{}}}, Main.workspace285.var"#1#2", Tuple{}})
   () @ Terminal.jl:77
7. with_logstate(::Function, ::Any) @ logging.jl:491
8. with_logger @ logging.jl:603 [inlined]
9. macro expansion @ Terminal.jl:76 [inlined]
10. macro expansion @ Suppressor.jl:165 [inlined]
11. macro expansion @ Terminal.jl:75 [inlined]
12. macro expansion @ Suppressor.jl:127 [inlined]
13. macro expansion @ Terminal.jl:74 [inlined]
14. macro expansion @ Suppressor.jl:206 [inlined]
15. var"#with_terminal#46"(::Base.Iterators.Pairs{Union{}, Union{}, Tuple{},
   NamedTuple{(), Tuple{}}}, ::typeof(PlutoUI.with_terminal),
   ::Function) @ Terminal.jl:73
16. with_terminal(::Function) @ Terminal.jl:72
17. top-level scope @ (Local: 1
```

```
• with_terminal() do
•   @time f(x, x̂)
• end
```

0.075599 seconds (62.15 k allocations: 5.979 MiB, 26.11% gc time, 71.91% compilation time)

```
• with_terminal() do
•   @time fastf(x, x̂)
• end
```

The error above is due to the missing vectorization of the function `f`. We can add this quite easily.

fastf (generic function with 1 method)

```

• # Can we get it down? Hide at first
• function fastf(x::AbstractVector{X}, y::AbstractVector{Y}) where {X, Y}
•   # We want a common type
•   T = promote_type(X, Y)
•   # Convert the input to a common type
•   x = convert.(T, x)
•   y = convert.(T, y)
•   # Preallocation of the result
•   res = Vector{T}(undef, length(x))
•   # Vectorize
•   # We assume length(x) == length(y) -> Errors are not caught!
•   # Additionally we add the @simd macro to instruct the compiler
•   # to parallelize
•   @simd for i in 1:length(x)
•       # If we want, we could also use the @fastmath macro here
•       # which does not improve our performance ( in this case )
•       res[i] = f(x[i], y[i])
•   end
•   # The return argument
•   return res
• end

```

fastf! (generic function with 1 method)

```

• # Mutating function and even faster because we assume the same type
• function fastf!(res::AbstractVector{T}, x::AbstractVector{T}, y::AbstractVector{T})
•   where {T}
•   # Vectorize
•   # We assume length(x) == length(y) -> Errors are not caught!
•   # Additionally we add the @simd macro to instruct the compiler
•   # to parallelize
•   @inbounds @simd for i in 1:length(x)
•       # If we want, we could also use the @fastmath macro here
•       # which does not improve our performance ( in this case )
•       res[i] = f(x[i], y[i])
•   end
•   # The return argument
•   return
• end

```

**MethodError: no method matching** `^(::Vector{Float64}, ::Int64)`

Closest candidates are:

`^(!Matched::Union{AbstractChar, AbstractString}, ::Integer)` at `strings/basic.jl:718`

`^(!Matched::Complex{var"#s79"} where var"#s79"<:AbstractFloat, ::Integer)` at `complex.jl:820`

`^(!Matched::Complex{var"#s79"} where var"#s79"<:Integer, ::Integer)` at `complex.jl:820`

...

1. `macro expansion` @ `none:0` [inlined]
2. `literal_pow` @ `none:0` [inlined]
3. `f(::Vector{Int64}, ::Vector{Float64})` @ **( Other: 2**
4. `macro expansion` @ `timing.jl:210` [inlined]
5. `(::Main.workspace285.var"#mybench#7") (::Function, ::Vector{Int64}, ::Vararg{Any, N} where N)` @ **( Local: 3**
6. `(::Main.workspace285.var"#5#6")()` @ **( Local: 5**
7. `(::PlutoUI.var"#49#52"{Base.Iterators.Pairs{Union{}, Union{}, Tuple{}, NamedTuple{(), Tuple{}}}, Main.workspace285.var"#5#6", Tuple{}})`  
`()` @ `Terminal.jl:77`
8. `with_logstate (::Function, ::Any)` @ `logging.jl:491`
9. `with_logger` @ `logging.jl:603` [inlined]
10. `macro expansion` @ `Terminal.jl:76` [inlined]
11. `macro expansion` @ `Suppressor.jl:165` [inlined]
12. `macro expansion` @ `Terminal.jl:75` [inlined]
13. `macro expansion` @ `Suppressor.jl:127` [inlined]
14. `macro expansion` @ `Terminal.jl:74` [inlined]
15. `macro expansion` @ `Suppressor.jl:206` [inlined]
16. `var"#with_terminal#46" (::Base.Iterators.Pairs{Union{}, Union{}, Tuple{}, NamedTuple{(), Tuple{}}}, ::typeof(PlutoUI.with_terminal), ::Function)` @ `Terminal.jl:73`
17. `with_terminal (::Function)` @ `Terminal.jl:72`
18. `top-level scope` @ **( Local: 1**

```

• with_terminal() do
•   # Benchmarking
•   mybench(f::Function, args...) = @time f(args...)
•
•   mybench(f, x, x̂)
•   mybench(fastf, x, x̂)
•
•   res = similar(x̂)
•   mybench(fastf!, res, convert.elttype(x̂), x), x̂)
• end

```

# Benchmarks



Microbenchmarks of Julia vs. different Languages as currently available [here](#)

**BUT** This is an old plot. Be careful.

► Just In Time (JIT) Compilation

## Some Performance Pitfalls

- Benchmarking
- Memory Layout
- Referecing

## Batteries included

Within the [LinearAlgebra](#) package - included in the standard library - we have BLAS and LAPACK at our fingertips.

10.0

```
• # Simple dot
• BLAS.dot(10, fill(1.0, 10), 1, fill(1.0, 20), 2)
```

[6, 9, 12]

```
• # a*x + y
• BLAS.axpy!(2, [1;2;3], [4;5;6])
```

(3×3 Matrix{Float64}:  
1.0 0.0 -2.16163e-15 0.744725 0.194975 1.78919  
3.97649e-16 0.0 1.0 -0.58645 1.86899 2.56491  
-2.68114e-16 1.0 1.07246e-15 0.523223 -0.726183 -0.207043

```
• # Solve A X = B via LU(A) and overwrites B with the solution
• begin
•     A = randn(3,3)
•     B = A * [1 0 0; 0 0 1; 0 1 0]
•     LAPACK.gesv!(A, B)
• end
```

# Classical Dispatch

## Functional

Global Namespace

$$f \in \mathcal{F}$$

Unique Functions

$$f(x_1, x_2, x_3, \dots)$$

Results in

```
z = 3+i4
r = 0.1
complex_real_add(z, r)
```

## Object Oriented

Local Namespace

$$f \in \mathcal{O}$$

Unique Functions in Namespace

$$o.f(x_1, x_2, x_3, \dots)$$

Can result in

```
z = 3+i4
r = 0.1
z + r
```



# Multiple Dispatch

Global namespace

$$f \in \mathcal{F}$$

With unique arguments

$$f : \mathcal{X}_1 \times \mathcal{X}_2 \cdots \mapsto \mathcal{Y}$$

```
z = 3+im
r = 0.1
z + r
```

► Is this useful?

► But why exactly is this useful?

## Dual Numbers

$$z = a + b\epsilon$$

with

$$\epsilon^2 = 0$$

```
• # Type definition
• struct DualNumber{T} <: Number
•     a::T
•     b::T
• end
```

```
• # Addition
• Base.:+(x::DualNumber{T}, y::DualNumber{T}) where T = DualNumber(x.a + y.a, x.b + y.b)
```

```
z = DualNumber(1.0, 0.2)
```

```
• z = DualNumber(1.0, 0.2)
```

```
DualNumber(100.0, 20.0)
```

```
• sum([z for i in 1:100])
```

```

• # Multiplication
• Base.:*(x::DualNumber, y::DualNumber) = DualNumber(
•     x.a*y.a, x.a*y.b+y.a*x.b
• )

```

```
DualNumber(1.0, 0.4)
```

```
• z*z # Works
```

```
DualNumber(1024.0, 5120.0)
```

```
• prod([(z^3 + z^2) for i in 1:10])
```

```

• begin
•     # Add multiplication
•     Base.:*(x::Number, y::DualNumber{T}) where T = DualNumber(convert(T, x)*y.a,
•         convert(T,x)*y.b)
•     Base.:*(y::DualNumber{T}, x::Number) where T = DualNumber(convert(T, x)*y.a,
•         convert(T,x)*y.b)
•     # Add subtraction
•     Base.:-(x::Number, y::DualNumber{T}) where T = DualNumber(convert(T,x)-y.a, -y.b)
•     Base.:-(y::DualNumber{T}, x::Number) where T = DualNumber(y.a-convert(T,x), y.b)
•     Base.:-(x::DualNumber, y::DualNumber) = x+(-1*y)
• end

```

```
f (generic function with 2 methods)
```

```
• f(x) = x^2 - 3*(x-2)^4
```

```
DualNumber(-2.0, 2.8000000000000003)
```

```
• f(z)
```

```
-24.966299999999993
```

```
• f(0.3)
```

```
-218
```

```
• f(5)
```


```
191.2137 - 372.13199999999995im
```

```
• f(0.3+im*3)
```

# Matchig Pursuit [4]

---

$$\min_{\Xi} \|\Xi\|_0, \quad \text{s.t.} \quad Y = \Psi(X)\Xi$$



## Pseudo-Code

- Compare each element of the normalized dictionary  $\Psi(X)$  to the signal  $Y$  via the inner product
- Use the largest resemblance as a coefficient  $\xi = \Xi_{i,j}$
- Subtract the recovered signal and repeat until converged

We will use a more sophisticated version **based on this paper** where our goal is to find the right support  $\Lambda$  of the coefficient vector.

*Note*

LinearAlgebra is already imported

gOMP (generic function with 4 methods)

```

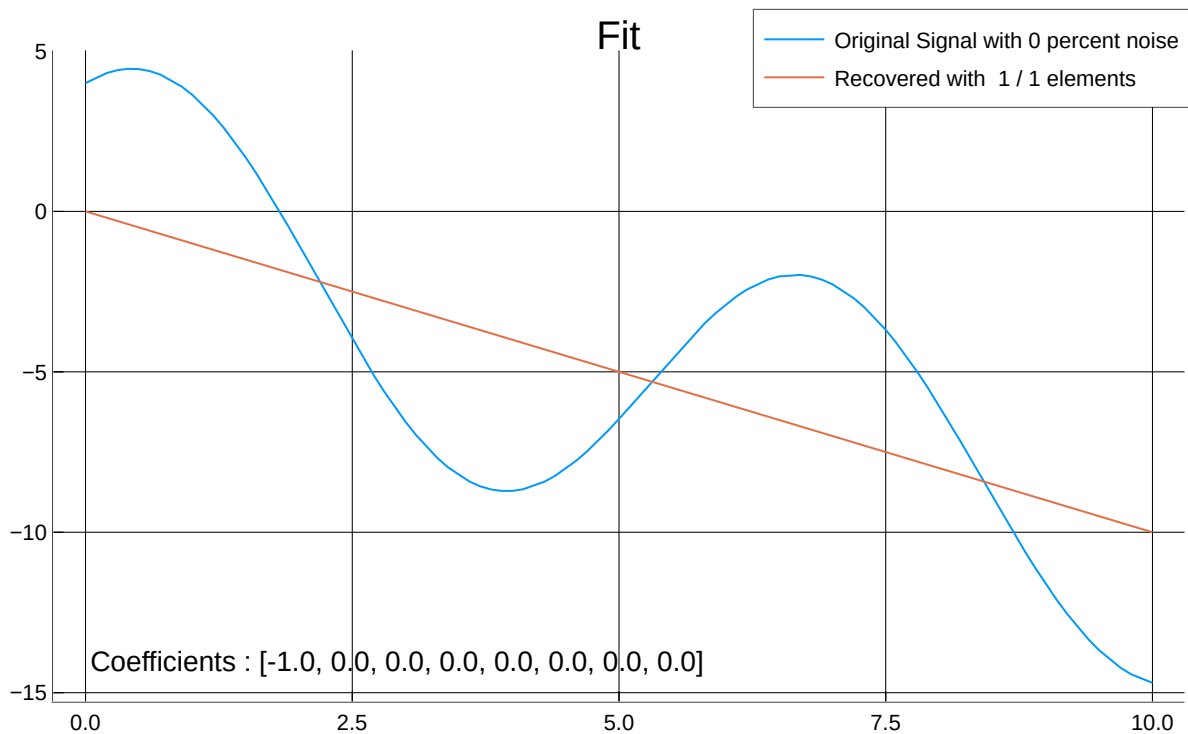
• # Define the algorithm
• function gOMP(y0::AbstractVector, Ψ::AbstractMatrix, K::Int = 2, S::Int = 1;
  max_iter::Int = 100, ε::Real = eps())
•   # Get the dimensions
•   m = length(y0)
•   m_psi, n = size(Ψ)
•   # Assert the dimensionality
•   @assert m == m_psi "Please provide consistent input sizes"
•   # Assert the selector
•   @assert S <= min(K, m/K) "S <= min(K, m/K)"
•
•   # Normalize
•   ψ = deepcopy(Ψ)
•   normalize!.(eachcol(ψ), 2)
•
•
•   # Preconditioning
•   P = ψ'pinv(ψ*ψ')
•   # New matrix
•   ψ = P*ψ
•   y = P*y0
•   # Iteration
•   iters = 0
•
•   # Support
•   Λ = zeros{Bool, n}
•   u = zeros{eltype(y), n}
•   r = y
•   amps = zeros{eltype(y), n}
•
•
•   # Find the magnitudes
•   for i in 1:max_iter
•     # Compute the similarities via magnitude
•     amps .= ψ'r
•     # Get the largest entry
•     idx = sortperm(abs.(amps), rev = true)[1:S]
•     # Update the support
•     Λ[idx] .= true
•     # Update the coefficients
•     u[Λ] .= (y' / ψ')[1, Λ]
•     # Update r
•     r .= y - ψ*u
•     # Convergence
•     if norm(r,2) < ε || sum(Λ) >= K
•       # Just for debug
•       #@debug "Early break after $i iterations with $(norm(r,2))"
•       break
•     end
•   end
•   # Last time to get the right coefficients
•   u[Λ] .= (y0' / ψ')[1, Λ]
•   return u
•
• end

```

```

• # Generate some test data
• begin
•     # Independent variable
•     t = 0.0:0.1:10.0;
•     # Signal
•     y = 3.0*sin.(t).*exp.(-t./50.0) + 4.0*cos.(t) - t;
•     # Dictionary
•     ψ = [
•         t t.^2 t.^3 sin.(t) ones(eltype(y), length(y)) cos.(t) exp.(t) sin.(t).*exp.
•         (-t./50.0)
•     ];
•     nothing
• end

```



Lets explore the data by adding a slider for the sparsity and a noise level.

Sparsity :

Noise :

gOMP (generic function with 4 methods)

```

• # What about parallism?
•
• # Simple
•
• function gOMP(Y::AbstractMatrix, Ψ::AbstractMatrix, args...; kwargs...)
•
•     # we know that the coefficients can be derived independent
•     A = zeros(size(Ψ, 2), size(Y, 1)) # Init
•     #Threads.@threads for i in 1:size(Y, 1)
•     for i in 1:size(Y, 1)
•         A[:, i] .= gOMP(Y[i, :], Ψ, args...; kwargs...)
•     end
•     return A
• end

```

```

Y = 100×101 Matrix{Float64}:
 3.9926  4.17834  4.3087  4.38064  ... -14.2115 -14.4429 -14.5998 -14.6769
 4.00903 4.20684  4.30776  4.3943   ... -14.2196 -14.4291 -14.5958 -14.6988
 4.00066 4.15286  4.30094  4.40755   ... -14.2102 -14.4186 -14.6069 -14.6678
 4.02103 4.17958  4.3141  4.4129   ... -14.2252 -14.4364 -14.5861 -14.6815
 4.01213 4.18151  4.33174  4.40018   ... -14.2218 -14.4331 -14.6001 -14.698
 4.00256 4.15751  4.30317  4.40488   ... -14.2306 -14.4041 -14.5932 -14.6867
 4.00827 4.16841  4.31768  4.40545   ... -14.2219 -14.4204 -14.5807 -14.6984
 ⋮
 4.01715 4.16797  4.29437  4.38201   ... -14.2401 -14.4049 -14.5691 -14.6966
 3.98759 4.18168  4.30862  4.38865   ... -14.2154 -14.4305 -14.5838 -14.6859
 3.98808 4.18097  4.31157  4.41082   ... -14.222  -14.4262 -14.5837 -14.6836
 3.99307 4.17842  4.32208  4.39345   ... -14.2258 -14.4091 -14.5824 -14.684
 4.02232 4.1989  4.32337  4.41797   ... -14.2194 -14.4256 -14.5855 -14.7017
 4.01931 4.1682  4.30764  4.39363   ... -14.2308 -14.4153 -14.5976 -14.6775

```

```

• Y = vcat([(y+1e-2*randn(size(y)))' for i in 1:100]...)

```

0.209139 seconds (165.72 k allocations: 90.200 MiB, 19.42% gc time, 47.75% compilation time)

```

• with_terminal() do
•     @time gOMP(Y, Ψ, 4)
• end

```

```

8×1 Matrix{Float64}:
-0.9938917657989252
-0.0017979565777127562
-1.033544637337606e-6
 0.003759089684817375
 0.0
 4.002800632607579
 0.0
 2.948770296238229

```

```

• sum(gOMP(Y, Ψ, 4), dims = 2)/100

```

# Ecosystem & Package Development

---

**JuliaHub** provides us with a nice, searchable database for all registered packages.

As an example for Package Development, we can have a look at **DataDrivenDiffEq.jl**.

And : **JuliaCon 2021** is around the corner!